

Introduction to statistics: Linear mixed models

Shravan Vasishth

Universität Potsdam
vasishth@uni-potsdam.de
<http://www.ling.uni-potsdam.de/~vasishth>

February 16, 2020

Summary

1. We know how to do simple t-tests.
2. We know how to fit simple linear models.
3. We saw that the paired t-test is identical to the varying intercepts linear mixed model.

Now we are ready to look at linear mixed models in detail.

Linear models

Returning to our SR/OR relative clause data from English (Grodner and Gibson, Expt 1). First we load the data as usual (not shown).

```
gge1crit<-read.table("data/grodnergibson05data.txt",  
                    header=TRUE)  
  
gge1crit$so<-ifelse(gge1crit$condition=="objgap",1,-1)  
  
dat<- gge1crit  
dat$logrt<-log(dat$rawRT)  
  
bysubj<-aggregate(logrt~subject+condition,  
                  mean,data=dat)
```

Linear models


The simple linear model (incorrect for these data):

```
summary(m0<-lm(logrt~so,dat))$coefficients
```

##		Estimate	Std. Error	t value	Pr(> t)
##	(Intercept)	5.883056	0.019052	308.7841	0.0000000
##	so	0.062017	0.019052	3.2551	0.0011907

Linear models

We can visualize the different responses of subjects (four subjects shown):



Linear models

Given these differences between subjects, you could fit a separate linear model for each subject, collect together the intercepts and slopes for each subject, and then check if the intercepts and slopes are significantly different from zero.

We will fit the model using log reading times because we want to make sure we satisfy model assumptions (e.g., normality of residuals).

Linear models

There is a function in the package `lme4` that computes separate linear models for each subject: `lmList`.

```
library(lme4)  
  
## Loading required package: Matrix  
  
lmlist.fm1<-lmList(logrt~so|subject,dat)
```

Linear models

Intercept and slope estimates for three subjects:

```
lm1$`1`$coefficients
```

```
## (Intercept)          so
```

```
##      5.769617      0.043515
```

```
lm1$`28`$coefficients
```

```
## (Intercept)          so
```

```
##      6.10021      0.44814
```

```
lm1$`37`$coefficients
```


```
## (Intercept)          so
```

```
##      6.61699      0.35537
```


Linear models

One can plot the individual lines for each subject, as well as the linear model m_0 's line (this shows how each subject deviates in intercept and slope from the model m_0 's intercept and slopes).

Linear models



Linear models

To find out if there is an effect of RC type, you can simply check whether the slopes of the individual subjects' fitted lines taken together are significantly different from zero.

Linear models

```
t.test(coef(lm1ist.fm1)[2])

##
##  One Sample t-test
##
## data:  coef(lm1ist.fm1)[2]
## t = 2.81, df = 41, p-value = 0.0076
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  0.017449 0.106585
## sample estimates:
## mean of x
##  0.062017
```

Linear models

The above test is exactly the same as the paired t-test and the varying intercepts linear mixed model **on aggregated data**:

```
t.test(logrt~condition,bysubj,paired=TRUE)$statistic
```

```
##          t
```

```
## 2.8102
```

```
## also compare with linear mixed model:
```

```
summary(lmer(logrt~condition+(1|subject),  
            bysubj))$coefficients[2,]
```

```
##      Estimate Std. Error    t value
```

```
## -0.124033    0.044137   -2.810207
```

Linear models

- ▶ The above `lmList` model we fit is called **repeated measures regression**. We now look at how to model unaggregated data using the linear mixed model.
- ▶ This model is now only of historical interest, and useful only for understanding the linear mixed model, which is the modern standard approach.

Linear mixed models

- ▶ The **linear mixed model** does something related to the above by-subject fits, but with some crucial twists, as we see below.
- ▶ In the model shown in the next slide, the statement $(1|\text{subject})$ adjusts the grand mean estimates of the intercept by a term (a number) for each subject.

Linear mixed models

Notice that we did not aggregate the data here.

```
m0.lmer<-lmer(logrt~so+(1|subject),dat)
```

Abbreviated output:

Random effects:

Groups	Name	Variance	Std.Dev.
subject	(Intercept)	0.09983	0.3160
Residual		0.14618	0.3823

Number of obs: 672, groups: subject, 42

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	5.88306	0.05094	115.497
so	0.06202	0.01475	4.205

Linear mixed models

One thing to notice is that the coefficients (intercept and slope) of the fixed effects of the above model are identical to those in the linear model `m0` above.

The varying intercepts for each subject can be viewed by typing:

```
ranef(m0.lmer)$subject[,1][1:10]
```

```
##      [1] -0.1039283  0.0771948 -0.2306209  0.2341978  0.0088
##      [7] -0.2055713 -0.1553708  0.0759436 -0.3643671
```


Visualizing random effects

Here is another way to summarize the adjustments to the grand mean intercept by subject. The error bars represent 95% confidence intervals.

```
library(lattice)
print(dotplot(ranef(m0.lmer, condVar=TRUE)))
```

Visualizing random effects

```
## $subject
```



Linear mixed models

The model `m0.lmer` above prints out the following type of linear model. `i` indexes subject, and `j` indexes items.

Once we know the subject id and the item id, we know which subject saw which condition:

```
subset(dat,subject==1 & item == 1)

##    subject item condition rawRT so  logrt
## 6         1     1    objgap   320  1 5.7683
```

$$y_{ij} = \beta_0 + u_{0i} + \beta_1 \times so_{ij} + \epsilon_{ij} \quad (1)$$

The **only** new thing here is the by-subject adjustment to the intercept.


Linear mixed models

- ▶ Note that these by-subject adjustments to the intercept u_{0i} are assumed by lmer to come from a normal distribution centered around 0:

$$u_{0i} \sim \text{Normal}(0, \sigma_{u0})$$

- ▶ The ordinary linear model m0 has one intercept β_0 for all subjects, whereas the linear mixed model with varying intercepts m0.lmer has a different intercept $(\beta_0 + u_{0i})$ for each subject i .
- ▶ We can visualize the adjustments for each subject to the intercepts as shown below.

Linear mixed models



Formal statement of varying intercepts linear mixed model

i indexes subjects, j items.

$$y_{ij} = \beta_0 + u_{0i} + (\beta_1 + u_{1i}) \times so_{ij} + \epsilon_{ij} \quad (2)$$

Variance components:

- ▶ $u_0 \sim Normal(0, \sigma_{u0})$
- ▶ $\epsilon \sim Normal(0, \sigma)$

Linear mixed models

Note that, unlike the figure associated with the `lmlist.fm1` model above, which also involves fitting separate models for each subject, the model `m0.lmer` assumes **different intercepts** for each subject **but the same slope**.

We can have `lmer` fit different intercepts AND slopes for each subject.

Linear mixed models

Varying intercepts and slopes by subject

We assume now that each subject's slope is also adjusted:

$$y_{ij} = \beta_0 + u_{0i} + (\beta_1 + u_{1i}) \times so_{ij} + \epsilon_{ij} \quad (3)$$

That is, we additionally assume that $u_{1i} \sim Normal(0, \sigma_{u1})$.

```
m1.lmer<-lmer(logrt~so+(1+so||subject),dat)
```

Random effects:

Groups	Name	Variance	Std.Dev.
subject	(Intercept)	0.1006	0.317
subject.1	so	0.0121	0.110
Residual		0.1336	0.365


Number of obs: 672, groups: subject, 42

Fixed effects:

Linear mixed models

These fits for each subject are visualized below (the red line shows the model with a single intercept and slope, i.e., our old model m_0):



varying intercepts and slopes for each subject



Linear mixed models

Comparing lmer model with varying intercepts model

Compare this model with the lmer.fm1 model we fitted earlier:




Visualizing random effects

```
print(dotplot(ranef(m1.lmer, condVar=TRUE)))
```

Visualizing random effects

```
## $subject
```



Formal statement of varying intercepts and varying slopes linear mixed model

i indexes subjects, j items.

$$y_{ij} = \beta_0 + u_{0i} + \beta_1 \times so_{ij} + \epsilon_{ij} \quad (4)$$

Variance components:

- ▶ $u_0 \sim \text{Normal}(0, \sigma_{u0})$
- ▶ $u_1 \sim \text{Normal}(0, \sigma_{u1})$
- ▶ $\epsilon \sim \text{Normal}(0, \sigma)$


Shrinkage in linear mixed models

- ▶ The estimate of the effect by participant is smaller than when we fit a separate linear model to the subject's data.
- ▶ This is called shrinkage in linear mixed models: the individual level estimates are shunk towards the mean slope.
- ▶ The less data we have from a given subject, the more the shrinkage.


- └ Linear mixed models
- └ Shrinkage in linear mixed models

Shrinkage in linear mixed models


Subject 28's estimates



Subject 36's estimates



Subject 37's estimates



Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

Let's randomly delete some data from one subject:

```
set.seed(4321)
## choose some data randomly to remove:
rand<-rbinom(1,n=16,prob=0.5)
```

Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

```
dat[which(dat$subject==37),]$rawRT

## [1] 770 536 686 578 457 487 2419 884 3365 233
## [15] 1081 971

dat$deletedRT<-dat$rawRT
dat[which(dat$subject==37),]$deletedRT<-
  ifelse(rand,NA,
         dat[which(dat$subject==37),]$rawRT)
```

Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

Now subject 37's estimates are going to be pretty wild:

```
subset(dat, subject==37)$deletedRT
```

```
## [1] 770 NA 686 578 NA NA NA NA 3365 233  
## [15] NA 971
```

Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

```
## original no pooling estimate:
lmList.fm1_old<-lmList(log(rawRT)~so|subject,dat)
coefs_old<-coef(lmList.fm1_old)
intercepts_old<-coefs_old[1]
colnames(intercepts_old)<-"intercept"
slopes_old<-coefs_old[2]
## subject 37's original estimates:
intercepts_old$intercept[37]

## [1] 6.617

slopes_old$so[37]

## [1] 0.35537
```

Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

```
## on deleted data:
lmList.fm1_deleted<-lmList(log(deletedRT)~so|subject,dat)
coefs<-coef(lmList.fm1_deleted)
intercepts<-coefs[1]
colnames(intercepts)<-"intercept"
slopes<-coefs[2]
## subject 37's new estimates on deleted data:
intercepts$intercept[37]

## [1] 6.6879


slopes$so[37]

## [1] 0.38843
```

Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

Subject 37's estimates



Shrinkage in linear mixed models

The effect of missing data on estimation in LMMs

- ▶ What we see here is that the estimates from the hierarchical model are barely affected by the missingness, but the estimates from the no-pooling model are heavily affected.
- ▶ This means that linear mixed models will give you more robust estimates (think Type M error!) compared to no pooling models.
- ▶ This is one reason why linear mixed models are such a big deal.

Crossed subjects and items in LMMs

Subjects and items are fully crossed:

```
head(xtabs(~subject+item,dat))
```

##		item															
##	subject	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
##		1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
##		2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
##		3	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
##		4	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
##		5	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
##		6	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

Linear mixed models

Linear mixed model with crossed subject and items random effects.

```
m2.lmer<-lmer(logrt~so+(1+so||subject)+(1+so||item),dat)
```

Linear mixed models

Random effects:


Groups	Name	Variance	Std.Dev.
subject	(Intercept)	0.10090	0.3177
subject.1	so	0.01224	0.1106
item	(Intercept)	0.00127	0.0356
item.1	so	0.00162	0.0402
Residual		0.13063	0.3614

Number of obs: 672, groups: subject, 42; item, 16

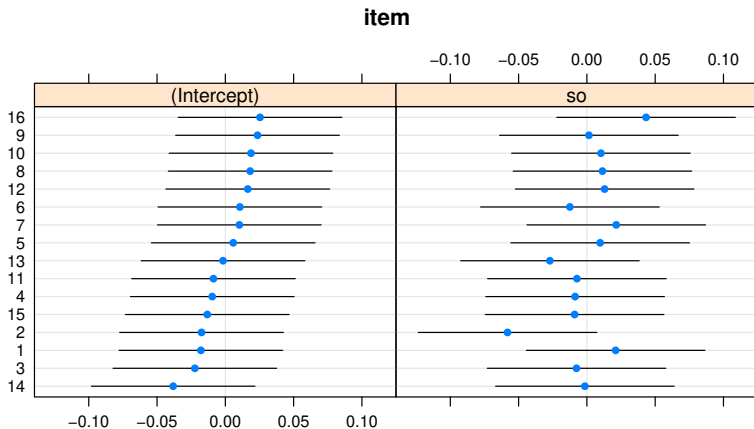
Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	5.8831	0.0517	113.72
so	0.0620	0.0242	2.56

Visualizing random effects



Visualizing random effects



Linear mixed models

Linear mixed model with crossed subject and items random effects.

```
m3.lmer<-lmer(logrt~so+(1+so|subject)+(1+so|item),  
              dat)
```

```
## boundary (singular) fit: see ?isSingular
```

Linear mixed models

Linear mixed model with crossed subject and items random effects.

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
subject	(Intercept)	0.10103	0.3178	
	so	0.01228	0.1108	0.58
item	(Intercept)	0.00172	0.0415	
	so	0.00196	0.0443	1.00 <= degenerate
Residual		0.12984	0.3603	

Number of obs: 672, groups: subject, 42; item, 16

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	5.8831	0.0520	113.09
so	0.0620	0.0247	2.51

Formal statement of varying intercepts and varying slopes linear mixed model with correlation

i indexes subjects, j items.


$$y_{ij} = \alpha + u_{0i} + w_{0j} + (\beta + u_{1i} + w_{1j}) * so_{ij} + \varepsilon_{ij} \quad (5)$$

where $\varepsilon_{ij} \sim Normal(0, \sigma)$ and

$$\Sigma_u = \begin{pmatrix} \sigma_{u0}^2 & \rho_u \sigma_{u0} \sigma_{u1} \\ \rho_u \sigma_{u0} \sigma_{u1} & \sigma_{u1}^2 \end{pmatrix} \quad \Sigma_w = \begin{pmatrix} \sigma_{w0}^2 & \rho_w \sigma_{w0} \sigma_{w1} \\ \rho_w \sigma_{w0} \sigma_{w1} & \sigma_{w1}^2 \end{pmatrix} \quad (6)$$

$$\begin{pmatrix} u_0 \\ u_1 \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \Sigma_u \right), \quad \begin{pmatrix} w_0 \\ w_1 \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \Sigma_w \right) \quad (7)$$

Visualizing random effects



Visualizing random effects

These are degenerate estimates

